

# Course Project 2

## Regular Expressions

CSE 30151 Spring 2023

Version of 2022-12-08

Due date 2023-03-10

In this project, you'll write a regular expression matcher similar to **grep**. This has three major steps: first, parse a regular expression into regular operations; second, execute the regular operations to create a NFA; third, run the NFA on input strings. Because we use a linear-time NFA recognition algorithm, our regular expression matcher will actually be much faster than one written using Perl or Python's regular expression engine. (Most implementations of **grep**, as well as Google RE2, are linear like ours.)

**You will need a correct solution for CP1 to complete this project.** If your CP1 doesn't work correctly (or you just weren't happy with it), you may use the official solution or another team's solution, as long as you properly cite your source.

## Getting started

To make sure your repository is up to date, please have one team member run the commands

```
git pull https://github.com/ND-CSE-30151/regexp-skeleton
git push
```

and then other team members should run `git pull`. The project repository should then include the following files (among others):


```
bin.{linux,darwin}/
  parse_re
  normalize_tree
  union_nfa
  concat_nfa
  star_nfa
  string_nfa
  re_to_nfa
  agrep
  compare_nfa
perl/
  bgrep.pl
tests/
  test-cp2.sh
cp2/
```

Please place the programs that you write into the `cp2/` subdirectory.

# 1 Parser

Note: This part and part 2.1 can be done in parallel.

In this first part, we'll write a parser for regular expressions. The goal is to input a regular expression like  $(ab|)^*$  and output a tree like



In plain text, we write this tree as


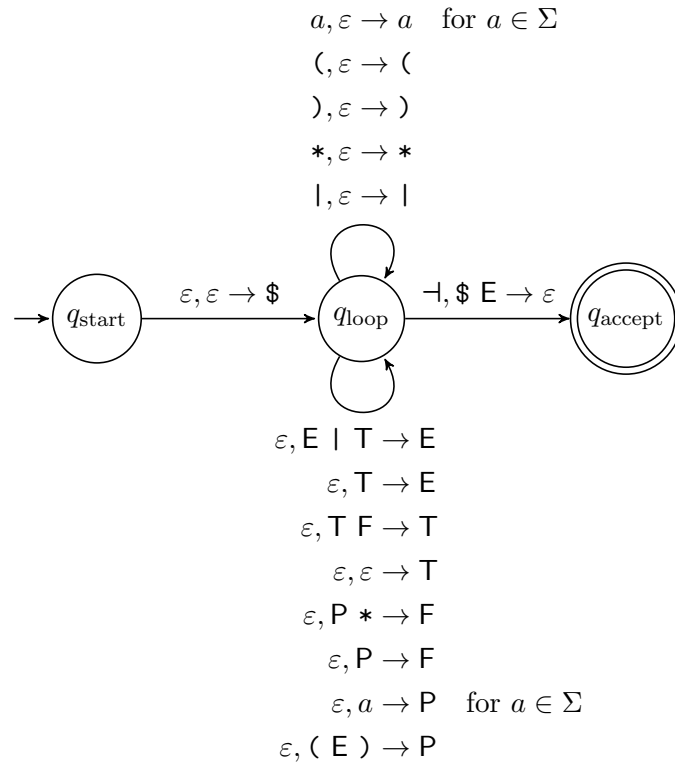
`(star (group (union (concat (symbol "a") (symbol "b")) (epsilon))))).`

Figure 1a shows a grammar for regular expressions. Here and below, let the terminal alphabet  $T$  be the set of all possible characters, and let  $\Sigma = T \setminus \{ (, ), *, |, \backslash \}$ . The nonterminal alphabet is  $V = \{E, T, F, P\}$ , with start symbol  $E$ .

A typical parser for a programming language is based on a PDA equivalent to the CFG. You've seen one way to convert a CFG to a PDA (Lemma 2.21), which is “top-down” in the sense that it starts with the start symbol ( $E$ ) and tries to rewrite it into the input string. There's another way to convert a CFG to a PDA, which works “bottom-up” in the sense that it starts with the string and tries to reduce it to  $E$ . It converts our grammar to the PDA in Figure 1b. Its input alphabet is  $T \cup \{\neg\}$ , where  $\neg$ , called the *endmarker*, must be appended to the input string. Its stack alphabet is  $\Gamma = V \cup T \cup \{\$\}$ . We are using a shorthand similar to that in the proof of Lemma 2.21, in which a transition can pop multiple symbols. However, we are writing stacks with the bottom to the left and the top to the right. For example, the transition  $\varepsilon, E \mid T \rightarrow E$  means “pop  $T$ , pop  $\mid$ , pop  $E$ , then push  $E$ .”

Table 3 shows an example run of this PDA. It starts in state  $q_{\text{start}}$ , pushes  $\$$ , and moves to state  $q_{\text{loop}}$ . Then, at each step, it follows a transition. The PDA is nondeterministic: for example, at step 12, three transitions seem to be possible:  $\mid, \varepsilon \rightarrow \mid$  and  $\varepsilon, T \rightarrow E$  and  $\varepsilon, \varepsilon \rightarrow T$ . To decide which transition to follow, we need some additional constraints, shown in Table 2. Each row in this table is a transition from  $q_{\text{loop}}$  to  $q_{\text{loop}}$ , and has a *below* set (which defaults to  $\Gamma$ ) and a *next* set (which defaults to  $T \cup \{\neg\}$ ). If a transition has label  $a, x \rightarrow y$ , where  $a \in T \cup \{\neg, \varepsilon\}$  and  $x \in \Gamma^*$ , it can be followed iff the remaining input string is  $abv\neg$  for some  $b \in \text{next}$  and  $v \in T^*$ , and the stack is  $uzx$  for some  $u \in \Gamma^*$  and  $z \in \text{below}$ . For example, at step 12, transition  $\mid, \varepsilon \rightarrow \mid$  is not possible because it has *below* =  $\{E\}$  but the stack symbol below  $\varepsilon$  is  $T$ ; similarly, transition  $\varepsilon, \varepsilon \rightarrow T$  is not possible because it has *below* =  $\{\$, (, \mid\}$ . But transition  $\varepsilon, T \rightarrow E$  is possible because it has *below* =  $\{\$, ($  and the stack symbol under  $T$  is  $($ , and it has *next* =  $\{\mid, ), \neg\}$  and the next input symbol is  $\mid$ . The constraints are designed so that at most one transition is possible.<sup>1</sup> If no transition is possible, the parser should print an error message and quit.

<sup>1</sup>If you want to learn how to derive such constraints yourself, we invite you to read Sipser, Section 2.4, or Prof. Thain's *Introduction to Compilers and Language Design*, Chapter 4, about how to build LR parsers. Our parser here is known as a (1,1)-BRC parser, which is less powerful but very easy to implement.


$$E \rightarrow T$$
$$T \rightarrow T F$$
$$T \rightarrow \varepsilon$$
$$F \rightarrow P *$$
$$F \rightarrow P$$
$$P \rightarrow a \quad \text{for } a \in \Sigma$$
$$P \rightarrow (E)$$


(a)

(b)

Figure 1: (a) A CFG for regular expressions, with start symbol  $E$ . Note that  $|$  is not being used here for writing two rules on one line. (b) A PDA converted from the CFG.

	below	next	tree
$a, \varepsilon \rightarrow a$ for $a \in \Sigma$	$\{T\}$		
$(, \varepsilon \rightarrow ($	$\{T\}$		
$), \varepsilon \rightarrow )$	$\{E\}$		
$ , \varepsilon \rightarrow  $	$\{E\}$		
$*, \varepsilon \rightarrow *$	$\{P\}$		
$\varepsilon, E \mid T \rightarrow E$		$\{ , ), \neg\}$	$(\text{union } \tau_1 \ \tau_2)$
$\varepsilon, T \rightarrow E$	$\{\$, (\}$	$\{ , ), \neg\}$	$\tau_1$
$\varepsilon, T \ F \rightarrow T$			$(\text{concat } \tau_1 \ \tau_2)$
$\varepsilon, \varepsilon \rightarrow T$	$\{\$, (,  \}$		$(\text{epsilon})$
$\varepsilon, P \ * \rightarrow F$			$(\text{star } \tau_1)$
$\varepsilon, P \rightarrow F$		$\Sigma \cup \{ (, ),  , \neg \}$	$\tau_1$
$\varepsilon, a \rightarrow P$ for $a \in \Sigma$			$(\text{symbol "a"})$
$\varepsilon, ( \ E \ ) \rightarrow P$			$(\text{group } \tau_1)$

Table 2: Additional information for transitions from  $q_{\text{loop}}$  to  $q_{\text{loop}}$  of the PDA in Figure 1b.

	stack	input	transition	pushed tree
1	\$	(ab )*¬	$\varepsilon, \varepsilon \rightarrow T$	(epsilon)
2	\$T	(ab )*¬	$(, \varepsilon \rightarrow ($	
3	\$T(	ab )*¬	$\varepsilon, \varepsilon \rightarrow T$	(epsilon)
4	\$T(T	ab )*¬	$a, \varepsilon \rightarrow a$	
5	\$T(Ta	b )*¬	$\varepsilon, a \rightarrow P$	(symbol "a")
6	\$T(TP	b )*¬	$\varepsilon, P \rightarrow F$	(symbol "a")
7	\$T(TF	b )*¬	$\varepsilon, TF \rightarrow T$	(symbol "a")
8	\$T(T	b )*¬	$b, \varepsilon \rightarrow b$	
9	\$T(Tb	)*¬	$\varepsilon, b \rightarrow P$	(symbol "b")
10	\$T(TP	)*¬	$\varepsilon, P \rightarrow F$	(symbol "b")
11	\$T(TF	)*¬	$\varepsilon, TF \rightarrow T$	(concat (symbol "a") (symbol "b"))
12	\$T(T	)*¬	$\varepsilon, T \rightarrow E$	(concat (symbol "a") (symbol "b"))
13	\$T(E	)*¬	$ , \varepsilon \rightarrow  $	
14	\$T(E	)*¬	$\varepsilon, \varepsilon \rightarrow T$	(epsilon)
15	\$T(E T	)*¬	$\varepsilon, E T \rightarrow E$	(union (concat (symbol "a") (symbol "b")) (epsilon))
16	\$T(E	)*¬	$), \varepsilon \rightarrow )$	
17	\$T(E)	*¬	$\varepsilon, (E) \rightarrow P$	(group (union (concat (symbol "a") (symbol "b")) (epsilon)))
18	\$TP	*¬	$*, \varepsilon \rightarrow *$	
19	\$TP*	¬	$\varepsilon, P* \rightarrow F$	(star (group (union (concat (symbol "a") (symbol "b")) (epsilon))))
20	\$TF	¬	$\varepsilon, TF \rightarrow T$	(star (group (union (concat (symbol "a") (symbol "b")) (epsilon))))
21	\$T	¬	$\varepsilon, T \rightarrow E$	(star (group (union (concat (symbol "a") (symbol "b")) (epsilon))))
22	\$E	¬	$\neg, \$E \rightarrow \varepsilon$	

Table 3: Example run of the parser. The stack is written with the bottom to the left and the top to the right.

In order to make the parser output a tree (not just an accept/reject decision), every transition that pushes a nonterminal  $Y$  attaches a subtree to  $Y$ . Table 2 shows, for each such transition, what the tree to be attached to  $Y$  is. The variables  $\tau_1$  and  $\tau_2$  stand for the trees attached to the first and second popped nonterminals. For example, in Table 3, at step 11, since transition  $\varepsilon, T F \rightarrow T$  has  $\text{tree} = (\text{concat } \tau_1 \tau_2)$ , we push  $T$  with a tree that has a root with label `concat` and two children: `(symbol "a")`, attached to the popped  $T$ , and `(symbol "b")`, associated with the popped  $F$ .

For efficiency's sake (in particular, if your Part 2.1 is recursive and you find that it exceeds the maximum recursion depth), you can apply simplifications like the following:

- `(concat (epsilon)  $\alpha$ )` becomes  $\alpha$ .
- `(union (union  $\alpha$   $\beta$ )  $\gamma$ )` becomes `(union  $\alpha$   $\beta$   $\gamma$ )`.
- `(concat (concat  $\alpha$   $\beta$ )  $\gamma$ )` becomes `(concat  $\alpha$   $\beta$   $\gamma$ )`.

For example, in Table 3, at step 7, the tree would have been `(concat (epsilon) (symbol "a"))`, but is simplified to just `(symbol "a")`. The automated tester simplifies your trees for you, so a correct parser will pass the tests with or without simplification.

Finally, at the end of the input string, if the stack is `$ E`, the parser moves to  $q_{\text{accept}}$  and outputs the tree associated with the popped  $E$ . That's the end! It will probably be helpful at this point to review the entire example run in Table 3.

Write code to run the parser described above on a regular expression and construct a tree from it. To test your parser, write a program with the following usage:

`parse_re regexp`

- *regexp*: a regular expression
- Output: string representation of the syntax tree for *regexp*

For example:

```
$ parse_re 'a'
(symbol "a")
$ parse_re ''
(epsilon)
$ parse_re '(a)'
(group (symbol "a"))
$ parse_re '()'
(group (epsilon))
$ parse_re 'a*'
(star (symbol "a"))
$ parse_re 'abc'
(concat (symbol "a") (symbol "b") (symbol "c"))
$ parse_re 'a|b|c'
(union (symbol "a") (symbol "b") (symbol "c"))
$ parse_re '||'
(union (epsilon) (epsilon) (epsilon))
```

Test your program by running `test-cp2.sh`.

## 2 Converter

### 2.1 Regular operations

Write a function that creates an NFA that accepts exactly one string, and a program to test it:

`string_nfa w`

- *w*: a string (possibly empty)
- Output: an NFA recognizing the language  $\{w\}$

Write functions that perform the three regular operations, using the constructions given in the book, and programs to test them:

`union_nfa  $M_1$   $M_2$`

- $M_1, M_2$ : NFAs
- Output: NFA recognizing language  $L(M_1) \cup L(M_2)$

`concat_nfa  $M_1$   $M_2$`

- $M_1, M_2$ : NFAs

- Output: NFA recognizing language  $L(M_1) \circ L(M_2)$

`star_nfa M`

- $M$ : an NFA
- Output: NFA recognizing language  $L(M)^*$

Test all of these programs by running `test-cp2.sh`.

## 2.2 Building the NFA

Write a function that converts (the syntax tree of) a regular expression to an NFA, by walking the tree bottom-up and using the operations implemented in Section 2.1. For now, the `group` operation should do nothing. Then combine your regular expression parser with this function into a test program:

`re_to_nfa regexp`

- *regexp*: Regular expression
- Output: NFA  $M$  equivalent to *regexp*

Test your program using `test-cp2.sh`.

## 3 Putting it together

Finally, combine your regular expression converter with your NFA simulator from CP1 to write a `grep` replacement, called `agrep` (for “automaton-based `grep`”):

`agrep regexp`

- *regexp*: regular expression
- Input: strings (one per line)
- Output: the input strings that match *regexp*

Note that unlike `grep`, the regular expression should match the entire line, not just part of the line. Test your program by running `tests/test-cp2.sh`.

The test script also tests the time complexity of `agrep`. This test is the same as in CP1, but now we can say a bit more about it. For various values of  $n$ , it creates the regular expression  $((aa|a)(a|aa))^n a^{4n}$  and tries to match it against the string  $a^{4n}$ , using our `agrep` and yours. For fun, we’ve provided a Perl implementation, called `bgrep.pl`, which you can try for comparison. (I ran out of patience and killed it.)

## Submission instructions

Your code should build and run on `studentnn.cse.nd.edu`. The automatic tester will clone your repository, `cd` into its root directory, run `make -C cp2`, and run `tests/test-cp2.sh`. You’re advised to try all of the above steps and ensure that all tests pass.

To submit your work, please push your repository to GitHub and then create a new release with tag version `cp2` (note that the tag version is not the same thing as the release title). If you are making a partial submission, then use a tag version of the form `cp2-123`, indicating which parts you’re submitting.

## Rubric

Part 1	
parsing	3
building syntax tree	3
parse_re	3
Part 2	
string_nfa	3
union_nfa	3
concat_nfa	3
star_nfa	3
re_to_nfa	3
Part 3 (agrep)	
correctness	3
time complexity	3
Total	30